

[• NUVERO.SPACE](#)

The AI infrastructure your **conversations** run on.

Conversational systems trained on how your business actually works, grounded in your real data, with a person always one step behind the ones a system can't answer alone.

4

LIVE DEPLOYMENTS

10+CHANNELS THE
ARCHITECTURE
SUPPORTS**0**ANSWERS THE MODEL
HAS TO GUESS AT**100%**TURNS LOGGED FOR
AUDIT

WHERE A SCRIPTED CHATBOT STOPS

A decision tree can't answer a question it was never scripted for.

No account context

The bot doesn't know who's asking or what state their account is actually in. Every reply is generic.

Ungrounded answers

It guesses at policy, eligibility and pricing instead of citing a real source. Wrong answers cost more trust than no answer.

Blind to operations


It can't read the record the question is actually about, so it falls back to "please contact support."

The result is handoffs, re-work, and a knowledge base nobody trusts. And the cost of the bot itself scales with seats, not with conversations actually resolved.

— ONE SYSTEM BETWEEN YOUR CHANNELS AND YOUR SYSTEMS OF RECORD

Your data stays yours. The layer just knows how to use it.

WhatsApp, web, Telegram, whatever your customers use today, all answered by one system that reads your real records and never has to be retrained to add a channel.



Your CRM, core systems and databases stay the system of record. The layer in front of them is owned by you.

— HOW IT STAYS ACCURATE

It may only state a fact that came back from a real lookup.

Every answer is checked against a versioned knowledge base and, where it matters, a live fetch against your own systems. If a fetch fails, the reply says so and points to the source. It does not fill the gap with something plausible-sounding.

11+50


KB CHAPTERS + Q&A
PAIRS INDEXED

Top-6

SOURCES RETRIEVED PER
ANSWER

0

FACTS INVENTED WHEN A
FETCH FAILS



Attendance and debar margins
pulled per request from the
university's own portal. Never
stored, never guessed.

— THE SELF-IMPROVING LOOP

A second model grades every answer. The knowledge base learns from what it finds.

Grading runs in the background on a stronger model, so it never slows the reply down. Findings land in a review queue; approved fixes go straight into the knowledge base the whole system draws from.




WHEN THE SYSTEM ISN'T SURE

It stops, hands the full conversation to a person, and waits.


- The system raises a flag and pauses itself for that one conversation, nothing else is affected.
- Whoever's on call gets the customer's name, the exact question, and a one-tap reply link.
- If nobody picks it up in 15 minutes, it un-pauses on its own so no one is left waiting.
- Every turn sits on an audit trail: the original message, what was retrieved, the final answer, the grade.

Data handling: conversation state lives in infrastructure you control, credentials are encrypted at rest, and nothing here is used to train a model for anyone else.


VITopia AI answers from a student's live university record.



Connect the portal once in chat, then ask for anything: today's schedule, attendance, fees.




Exam dates and seat number, fetched fresh from the portal at the moment of asking.




The same assistant, same brain, on Telegram, for students who never open the web app.

- About 25 live university-portal modules answered on request: attendance, marks, CGPA, exam seat, hostel room, pending assignments, fee receipts.
- When the portal throws a CAPTCHA, the assistant solves it itself by reading the image; a human only types one after four failed attempts.
- One brain across web and Telegram. Nothing personal is stored, and a failed fetch is labelled unverified rather than guessed at.


A support assistant for members 50+, in the language they actually write in.



A Hinglish question, answered in Hindi, in the same thread.



It recalls the member's previous question from a rolling conversation summary.



Walking a 50+ member through an app install, step by step, on WhatsApp.

- Replies in Hindi, Hinglish or English, whichever the member actually wrote in.
- Indic-aware moderation separates conversational filler from real frustration before a model ever runs, so tone is read correctly.
- Fully serverless: no VM, no Redis, scales to zero between conversations, backed by 28 automated pipeline tests.

— EVERY CHANNEL, ONE SYSTEM

Add a channel without retraining anything. The pipeline doesn't change.

WhatsApp Business

Vernacular support, in the thread.

Web app

Authenticated from the first message.

In-product SDK

iOS and Android, embedded in your app.

Telegram

Lightweight, no account needed.

Microsoft Teams

Agent-assist for support staff.

Slack

Internal helpdesk, in the channel.

Instagram DM

Same brain, where leads already are.

Messenger

Facebook Page inbox, answered live.

SMS / voice IVR

For customers offline apps miss.

Email

Drafts a grounded reply, not a form letter.

Same retrieval, same guardrails, same memory, every time. Channel code is the only thing that's different.

 ILLUSTRATIVE · MICRO-LENDING AND PERSONAL LOANS

The same architecture, pointed at a borrower's journey.

Eligibility guidance

Pre-qualification questions, a document checklist, an indicative rate, all inside chat.

KYC assistance

Walks the applicant through ID and selfie capture, recovers drop-offs over WhatsApp.

Repayment

Due-date reminders, balance questions, a payment link generated in the thread.

Vernacular support

The same language-matching GoHappy Club already runs in production today.

This pattern is drawn directly from the two deployments above. It has not yet been built for a lender, this is what the same infrastructure would look like pointed at that workflow.

ECONOMICS

Most messages never reach a large model.

- Junk and off-topic messages are dropped before any model is ever called.
- One rewrite, one retrieval, one generation per turn. The larger model only grades, and it does that in the background.
- Fully serverless: you pay for conversations that happen, not for capacity sitting idle.

PROOF OF CONCEPT • FOUR WEEKS

Scope the intents and data access. Build the pipeline against your stack. Ship one channel with escalation live. Run it on real traffic and review the numbers together.

MEASURED ON

Containment rate**CSAT****Cost / conversation****Escalation quality**

THE NEXT STEP

The future of support isn't more scripts.

It's a system that already knows your business, grounded in your real data, with a person always one step behind the questions it can't answer alone.

LET'S SCOPE YOUR PROOF OF CONCEPT

**nuvero.space**

Nuvero AI – the intelligence layer
underneath your conversations.